

# TP7

## Introduction:

### Structure du projet:

```
tp-dvc-mlflow/
|
|   ├── params.yaml
|   ├── dvc.yaml
|   ├── pyproject.toml
|   └── src/
|       ├── data/
|       |   ├── download.py
|       |   ├── preprocess.py
|       |   └── split.py
|       ├── train/
|       |   └── train.py
|       └── infer/
|           └── predict.py
|
|   └── data/
|       ├── raw/
|       ├── processed/
|       └── splits/
|
└── models/
└── reports/
└── mlruns/
```

## Dataset utilisé

Adult Income Dataset (UCI Machine Learning Repository)

Problème de classification binaire : prédire si le revenu annuel est >50K ou <=50K.

Sources :

<https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data>  
<https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.names>

Parameters globales params.yaml:

```
seed: 42

dataset:
    url: "https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data"
    raw_filename: "adult.csv"

preprocess:
    drop_missing: true
    drop_duplicates: true
    normalize_numeric: true
    categorical_encoding: "onehot"

split:
    train_ratio: 0.7
    val_ratio: 0.15
    test_ratio: 0.15
    stratify: true
```

Pour Lecture des paramètres YAML:

```
import yaml

with open("params.yaml", "r") as f:
    params = yaml.safe_load(f)

seed = params["seed"]
train_ratio = params["split"]["train_ratio"]
```

## Étape 1 – Téléchargement des données

Script : src/data/download.py

- Télécharger le dataset depuis l'URL définie dans params.yaml
- Sauvegarder le fichier CSV dans data/raw/

code :

```
import pandas as pd

url = params["dataset"]["url"]
columns = [
    "age", "workclass", "fnlwgt", "education", "education_num",
    "marital_status", "occupation", "relationship", "race", "sex",
    "capital_gain", "capital_loss", "hours_per_week",
    "native_country", "income"
]

df = pd.read_csv(url, header=None, names=columns)
df.to_csv("data/raw/adult.csv", index=False)
```

Le fichier brut doit être suivi par DVC.

```
dvc add data/raw/adult.csv
git add data/raw/adult.csv.dvc
Git commit -m 'add raw dataset'
```

## Étape 2 – Prétraitement et split

Script : src/data/[download.py](#)

**Transformations:**

- Remplacement de ? par NaN
- Suppression des doublons
- Suppression ou conservation des lignes incomplètes
- Normalisation des colonnes numériques

```
from __future__ import annotations
from pathlib import Path
```

```
import json
import yaml
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler

def load_params(path: str = "params.yaml") -> dict:
    with open(path, "r", encoding="utf-8") as f:
        return yaml.safe_load(f)
params = load_params()

raw_path = Path("data/raw") / params["dataset"]["raw_filename"]
out_dir = Path("data/processed")
out_dir.mkdir(parents=True, exist_ok=True)
out_path = out_dir / "adult_clean.csv"

df = pd.read_csv(raw_path)

df.replace("?", np.nan, inplace=True)

if params["preprocess"]["drop_duplicates"]:
    df = df.drop_duplicates()

if params["preprocess"]["drop_missing"]:
    df = df.dropna()

from sklearn.preprocessing import StandardScaler

num_cols = df.select_dtypes(include="number").columns
scaler = StandardScaler()
df[num_cols] = scaler.fit_transform(df[num_cols])

df.columns = [c.strip().lower().replace(" ", "_") for c in df.columns]

df.to_csv(out_path, index=False)
```

Script : [src/data/split.py](#)

```
params = load_params()
```

```

seed = int(params["seed"])
split_cfg = params["split"]
train_ratio = float(split_cfg["train_ratio"])
val_ratio = float(split_cfg["val_ratio"])
test_ratio = float(split_cfg["test_ratio"])
stratify = bool(split_cfg.get("stratify", True))

total = train_ratio + val_ratio + test_ratio
if abs(total - 1.0) > 1e-9:
    raise ValueError(f"Split ratios must sum to 1.0, got {total}")

in_path = Path("data/processed") / "adult_clean.csv"
df = pd.read_csv(in_path)

if "income" not in df.columns:
    raise ValueError("Target column 'income' not found. Check preprocessing output.")

X = df.drop(columns=["income"])
y = df["income"]

X_train, X_temp, y_train, y_temp = train_test_split(
    X,
    y,
    test_size=1.0 - train_ratio,
    random_state=seed,
    stratify=y if stratify else None,
)

# Split temp into val and test
val_ratio_adjusted = val_ratio / (val_ratio + test_ratio)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp,
    y_temp,
    test_size=1.0 - val_ratio_adjusted,
    random_state=seed,
    stratify=y_temp if stratify else None,
)

out_dir = Path("data/splits")
out_dir.mkdir(parents=True, exist_ok=True)

```

```

train_df = X_train.copy()
train_df["income"] = y_train.values
val_df = X_val.copy()
val_df["income"] = y_val.values
test_df = X_test.copy()
test_df["income"] = y_test.values

train_path = out_dir / "train.csv"
val_path = out_dir / "val.csv"
test_path = out_dir / "test.csv"

train_df.to_csv(train_path, index=False)
val_df.to_csv(val_path, index=False)
test_df.to_csv(test_path, index=False)

def dist(series: pd.Series) -> dict:
    vc = series.value_counts(normalize=True)
    return {str(k): float(v) for k, v in vc.items()}

print(f"[split] Input: {in_path} shape={df.shape}")
print(f"[split] seed={seed} stratify={stratify}")
print(f"[split] train={len(train_df)} val={len(val_df)} test={len(test_df)}")
print(f"[split] train class dist: {dist(train_df['income'])}")
print(f"[split] val class dist: {dist(val_df['income'])}")
print(f"[split] test class dist: {dist(test_df['income'])}")
print(f"[split] Saved: {train_path}, {val_path}, {test_path}")

if __name__ == "__main__":
    main()

```

## Étape 3 – Train

Script : [src/train/train.py](#)

```
from __future__ import annotations
```

```
from pathlib import Path
import json
import yaml
import pandas as pd

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

import joblib
import mlflow
import mlflow.sklearn

def load_params(path: str = "params.yaml") -> dict:
    with open(path, "r", encoding="utf-8") as f:
        return yaml.safe_load(f)

def main() -> None:
    params = load_params()
    seed = int(params["seed"])

    train_path = Path("data/splits/train.csv")
    test_path = Path("data/splits/test.csv")

    df_train = pd.read_csv(train_path)
    df_test = pd.read_csv(test_path)

    target = "income"
    X_train, y_train = df_train.drop(columns=[target]), df_train[target]
    X_test, y_test = df_test.drop(columns=[target]), df_test[target]

    cat_cols = X_train.select_dtypes(include=["object"]).columns.tolist()
    num_cols = [c for c in X_train.columns if c not in cat_cols]

    preprocessor = ColumnTransformer(
        transformers=[
```

```
("cat", OneHotEncoder(handle_unknown="ignore"), cat_cols),
("num", "passthrough", num_cols),
]
)

clf = Pipeline(
    steps=[
        ("prep", preprocessor),
        ("model", LogisticRegression(max_iter=1000, random_state=seed)),
    ]
)

mlflow.set_tracking_uri("file:/mlruns")
mlflow.set_experiment("adult-income-dvc-mlflow")

Path("models").mkdir(parents=True, exist_ok=True)
Path("reports").mkdir(parents=True, exist_ok=True)

with mlflow.start_run():
    mlflow.log_param("seed", seed)

    clf.fit(X_train, y_train)
    preds = clf.predict(X_test)
    acc = accuracy_score(y_test, preds)

    mlflow.log_metric("test_accuracy", float(acc))

    model_path = Path("models/model.joblib")
    joblib.dump(clf, model_path)
    mlflow.sklearn.log_model(clf, "model")

    metrics_path = Path("reports/metrics.json")
    with open(metrics_path, "w", encoding="utf-8") as f:
        json.dump({"test_accuracy": float(acc)}, f, indent=2)
    mlflow.log_artifact(str(metrics_path))

    print("[train] test_accuracy =", float(acc))
    print("[train] saved model:", model_path)
    print("[train] saved metrics:", metrics_path)
```

```
if __name__ == "__main__":
    main()
```

Avant d'exécuter ce script lancer mlflow dans le même répertoire:

```
> [poetry run] mlflow ui --backend-store-uri ./mlruns --port 5000
```

## Étape 4 – dvc pipeline

Script : **dvc.yaml**

```
stages:
  download:
    cmd: python src/data/download.py
    outs:
      - data/raw/adult.csv

  preprocess:
    cmd: python src/data/preprocess.py
    deps:
      - data/raw/adult.csv
    params:
      - preprocess
    outs:
      - data/processed/adult_clean.csv

  split:
    cmd: python src/data/split.py
    deps:
      - data/processed/adult_clean.csv
    params:
      - seed
      - split
    outs:
      - data/splits/train.csv
      - data/splits/val.csv
      - data/splits/test.csv
```

Exécuter Dvc repro pour lancer tous les etapes et entrainer

**dvc remove .\data\raw\adult.csv.dvc**

**dvc repro**